




Verifying Message-Passing Systems

Damien Zufferey

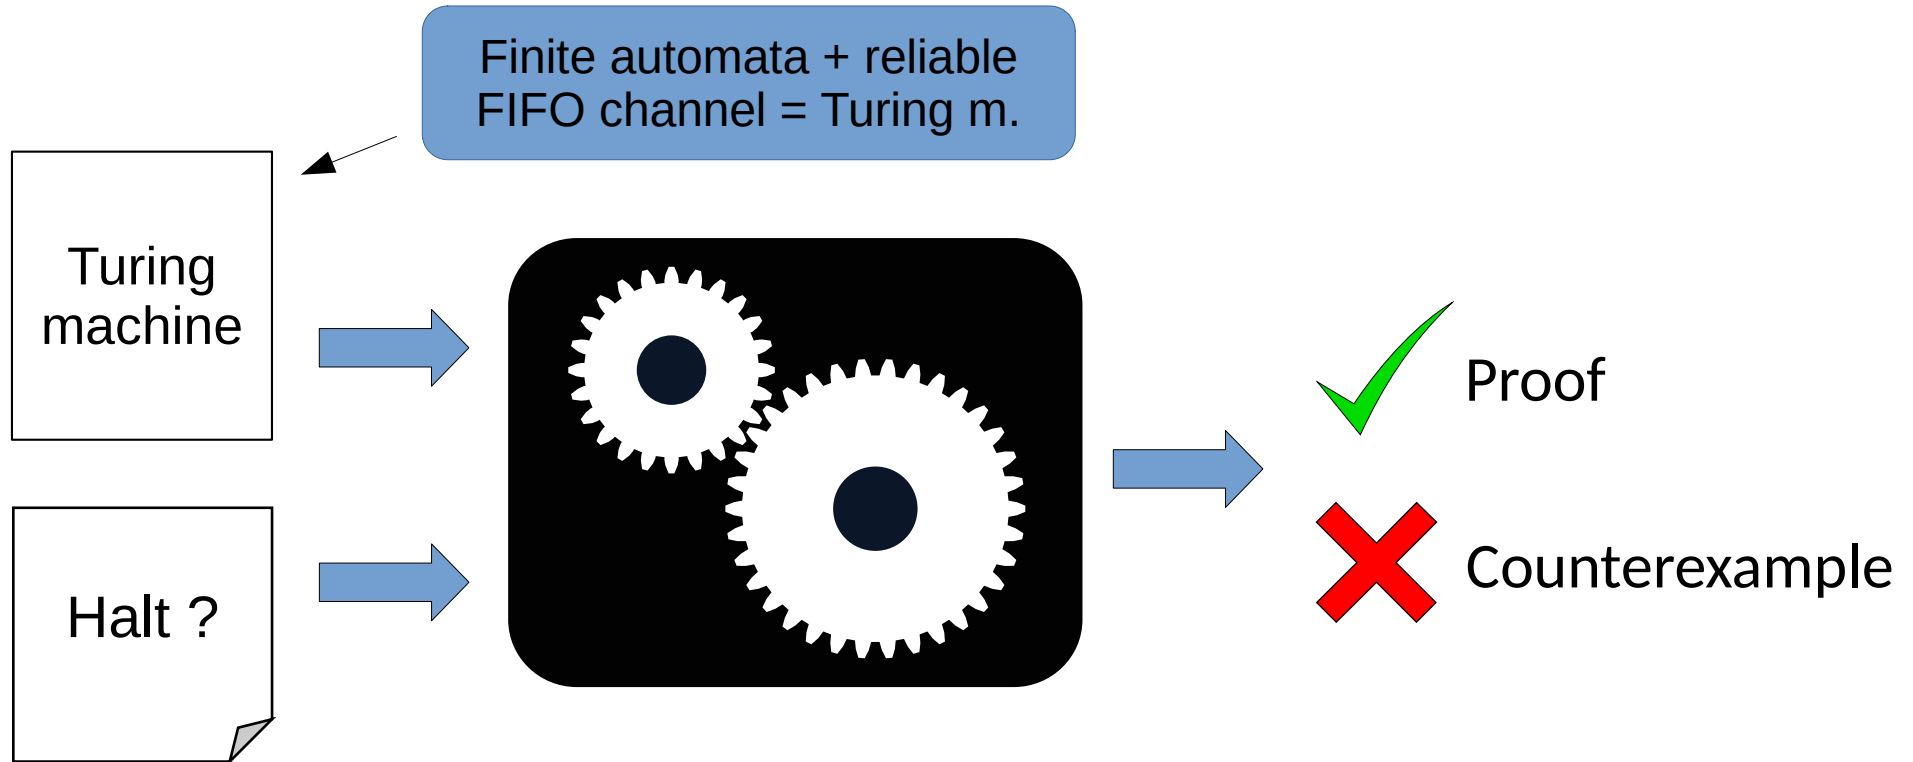
MPI-SWS

Dagstuhl 2019
28.10.2019

Automated Software Verification



Unfortunately ...



Problems Do Not Go Away Because They are Hard...

... but they still make life miserable.

Maybe the problem is ill-posed:

- Full automation is too much (ask the user for help)
- Try an easier problem (focus on null pointers)
- Programming languages are too powerful
- ...






Using PL to Make Verification Easier

Automated verification requires breaking a problem in small enough parts so the verification becomes “simple.”

Try to build a (de)composition strategy that the verification can use into the programming abstraction.

My research focus on doing that for message passing programs.

Some Earlier Work




P(rotocol) [PLDI 13]

- Communicating state-machines
- Decomposition: control and data
- Applied to device drivers in Windows


PSync (partial synchrony) [POPL 2016]

- Communication-closed rounds
- Decomposition: “time”
- Applied to fault-tolerant distributed algorithms




More Recent: Cyber-Physical Systems (CPS)

CPS = Communication + Computation + Control



- PGCD (Geometry, Concurrency, and Dynamics)
[ICCPS 19, ECOOP 19]
 - Decomposition: communication and local perspective
 - Application: modular robotics


PSync: A Partially Synchronous Language for Fault-tolerant Distributed Algorithms




Collaboration with Cezara Drăgoi,
Thomas A. Henzinger, Josef
Widder, Helmut Veith

[POPL 16, SNAPL 15, VMCAI 14] and
ongoing work

Why Care About Failure?



The Paxos Algorithm [Lamport 98]



Implementing Paxos: from ~50 lines of pseudo code to >500 LoC. What goes wrong?





Implementation Challenges


- Detecting failure (and the impossibility to get it right)
- Messages (side-effects) are untyped, have no scope, etc.
- Control-flow inversion (losing the program structure)

When Processes Fail?

Waiting for a message




Asynchrony



Wait longer can help but how to avoid blocking?

Asynchronous models are not good at dealing with faults.

Fault



[FLP 85] Solving consensus requires time.

Communication is a Side Effect

```
Type object;  
Byte[] buffer =  
    serialize(obj);  
send(channel, buffer);
```

011001010101...

```
buffer = recv(channel);  
Type1 object1 =  
    deserialize1(buffer);  
...  
buffer = recv(channel);  
Type2 object2 =  
    deserialize2(buffer);
```

Up to the programmer to:

- interpret the bytes moving over the network,
- know which receive corresponds to which send.

Control-flow Inversion

Protocol structure replaced by dispatch:

Protocol:
(1)Msg A
(2)Msg B







```
var state = 1

while (true) {
  on receive {
    case Msg A =>
      if (state == 1) ...
      else if (state == 2) ...
    case Msg B =>
      if (state == 1) ...
      else if (state == 2) ...
  }
}
```

normal case
message duplicated
message dropped
normal case

Verification Perspective

- Asynchrony
- Channels
- Faults
- Parametric systems



Interleavings → combinatorial explosion

reliable FIFO → undecidable

another combinatorial explosion

usually undecidable




PSync

- Goals: **make things simpler** (to understand and verify)
 - The code preserves the algorithm structure
 - Code complexity on par with pseudo-code
- How: **communication-closed rounds**
 - Syntactic scope to the messages
 - Provides logical time
 - Gives the *illusion* of synchrony

Communication-closed Rounds

[Elrad & Francez 82]




Paxos organized in communication-closed rounds.
No message cross the boundaries between rounds.

How Does It Model a Real System?


Idea: model faults/asynchrony as an adversarial environment [Gafni 98]
Project all the “faults” on the messages

Lockstep semantics:




Indistinguishable

Asynchronous execution:



Local views are preserved.


Round Example



```
new Round[(Int,Time)]{  
  
  def send(): Map[ProcessID, (Int,Time)] =  
    Map(coord -> (x, ts))  
  
  def update(mailbox: Map[ProcessID, (Int,Time)]) {  
    if (id == coord && mailbox.size > n/2) {  
      vote = valueWithMaxTS(mailbox)  
      commit = true  
    }  
  }  
}
```


Benefits for the Verification

- Asynchrony




No interleaving

- Channels




Buffer of size 1

- Faults



Project all faults on messages
(less nondeterminism)

- Parametric systems



Decidable with further restrictions
[Marić et al. CAV 17]

Problem: Performance and User-Control

Old version, runtime is a blackbox. No control:

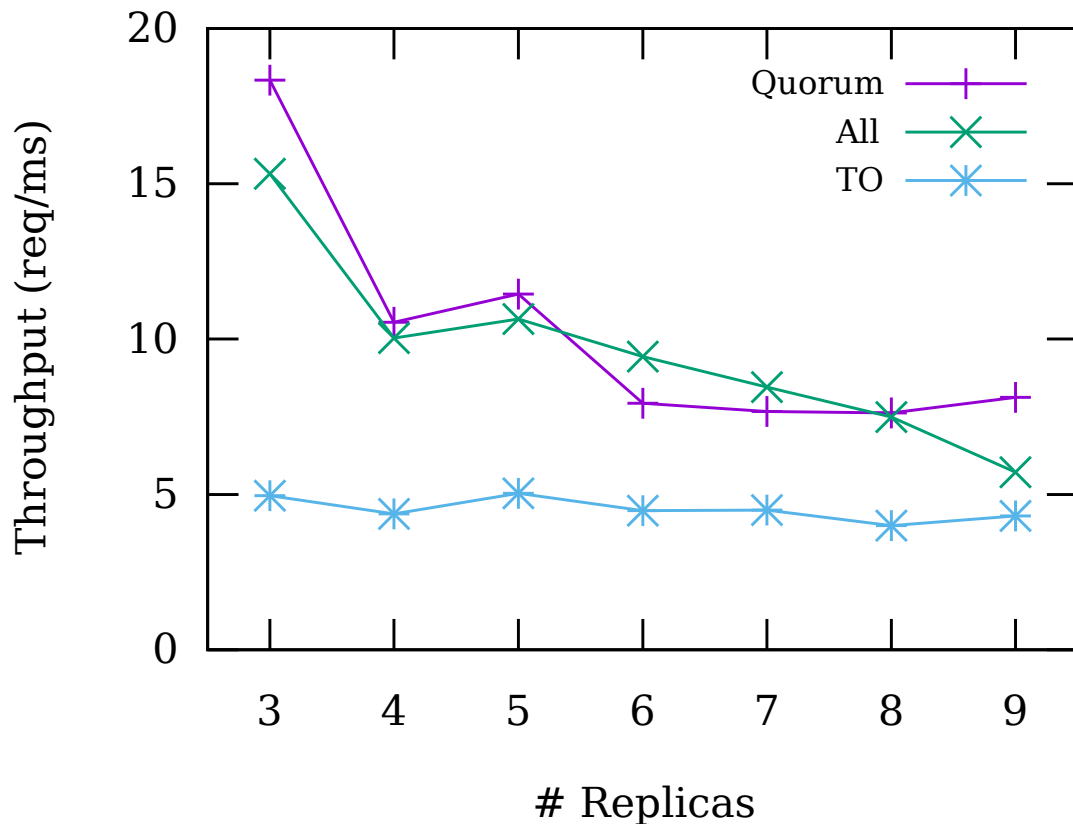
```
new Round[(Int,Time)]{  
  
  def send(): Map[ProcessID, (Int,Time)] =  
    Map(coord -> (x, ts))  
  
  def update(mailbox: Map[ProcessID, (Int,Time)]) {  
    if (id == coord && mailbox.size > n/2) {  
      ...  
    }  
  }  
}
```

Problem: Performance and User-Control

New version, the program can give hints to the runtime:

```
new Round[(Int,Time)]{  
  var nMsg = 0  
  def init =  
    if (id == coord) Progress.timeout( timeout )  
    else Progress.goAhead  
  def send(): Map[ID, (Int,Time)] =  
    Map(coord -> (x, ts))  
  def receive(sender: ID, payload: (Int, Time)) = {  
    nMsg += 1  
    ...  
    if (nMsg > n/2) Progress.goAhead  
    else Progress.unchanged  
  }  
}
```

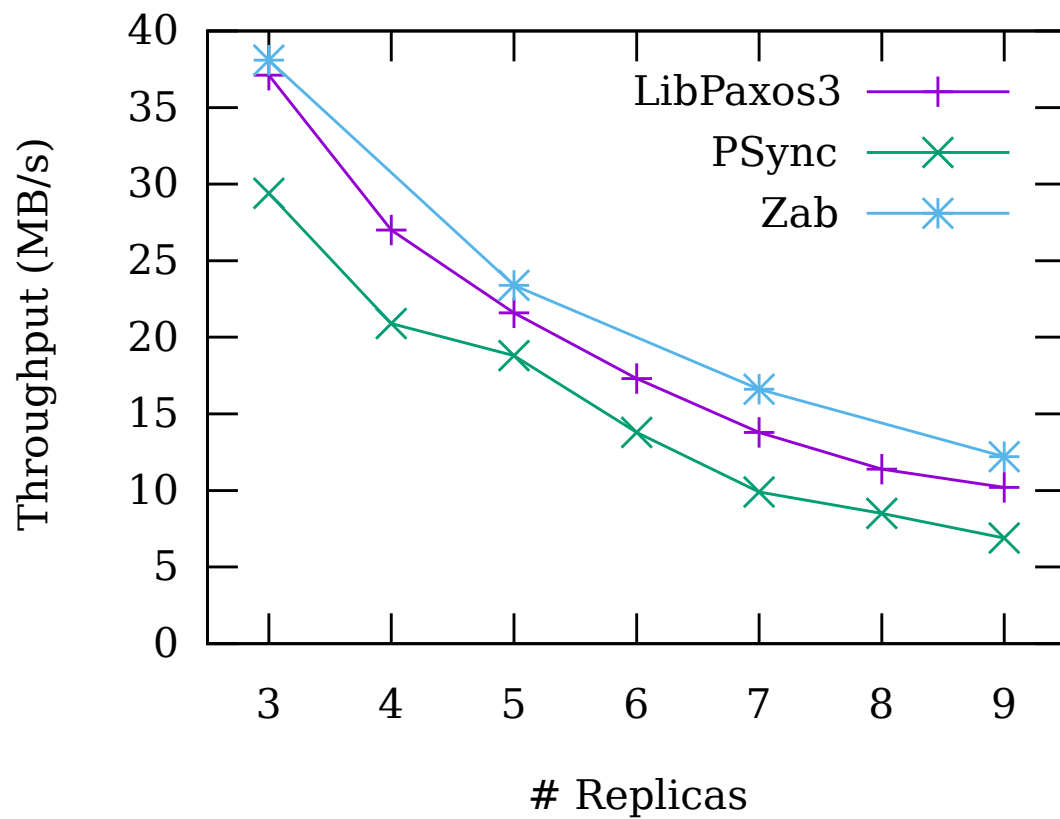
Results



Same algorithm but with different progress conditions.


Timeout (TO) is the best the old version can do will providing liveness guarantees.

Results



Current Work: Cyber-Physical Systems (CPS)

CPS = Communication + Computation + Control




- PGCD (Geometry, Concurrency, and Dynamics)
[ICCPS 19, ECOOP 19]
 - Decomposition: communication and local perspective
 - Application: modular robotics


Example: Handover



Example: Handover (Meet)




Example: Handover (Grab)




Example: Handover (Fold)




Example: Handover (Done)



How Many Robots?





How Many Robots?




Cart & Arm:
Two robots attached together that
act as “one” robot (communication)


How Many Robots?



Sequence of Actions



Arm's Code



```
while true do
  receive (Cart, idle)
  grab(loc) ⇒
    grip(loc);
    send(cart, ok)
  fold ⇒
    move(origin);
    send(cart, ok)
  done ⇒
    break
```

Robotic Program(mer)s as Target

- Message-Passing Abstraction
 - Robotic Operating System (ROS): publish-subscribe model
 - ROS is a de facto standard in academia (used for prototyping in industry)
- Modular robots
 - Software: logically separate units
 - Hardware: in a shared world (physical coupling)
- Different “kinds” of code
 - Planning: synchronize with other robots, decide what to do next, ...
 - Control: following trajectories, actuation, sensing, ...

Modularity for Components

Cart || Arm




Traditional parallel composition
lacks information.

Cart $\triangleleft_{\theta, M}$ Arm



Frame shift (relative position):
Arm's position depends on Cart

Dynamic coupling:
Cart's speed depends
on arm's weight.

Oriented operator:
- cart is the parent
- arm is the child




Making Composition Transparent




Goal: Verification

Global Spec.
 $\varphi(\text{Cart}, \text{Arm}, \text{Carrier})$



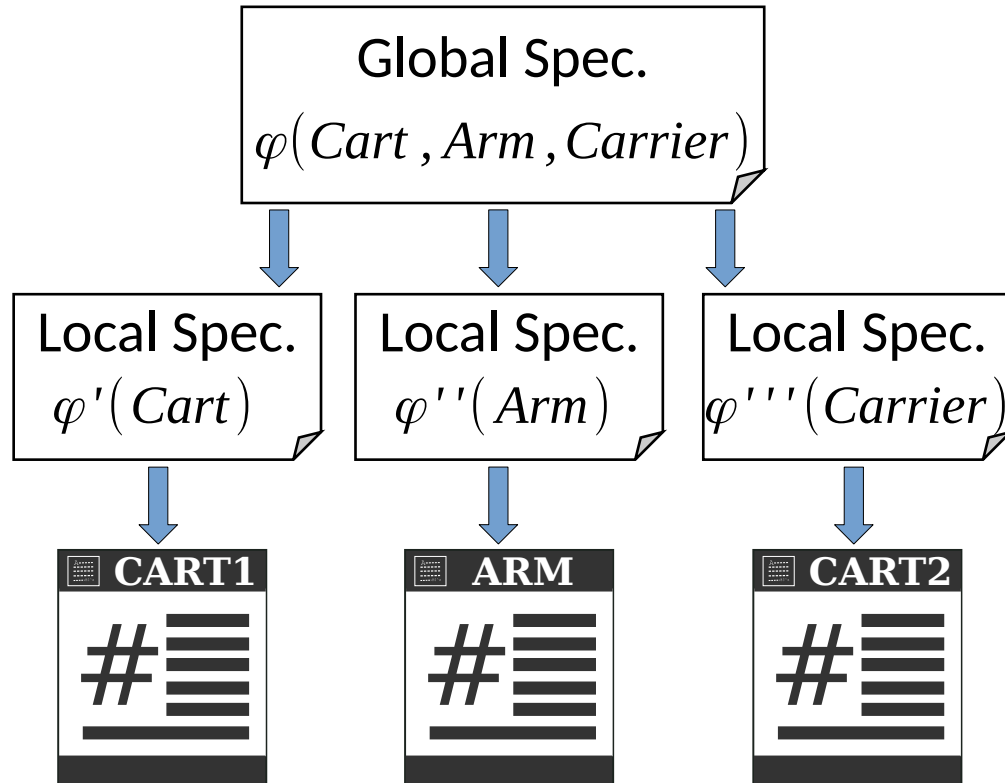
refines



The product is expensive.
(state-space explosion!)

Assume-Guarantee helps
for bottom-up composition.

Decomposing with Sessions Typs



Motion session type (Global)



Projections

Motion session type (Local)



Typing

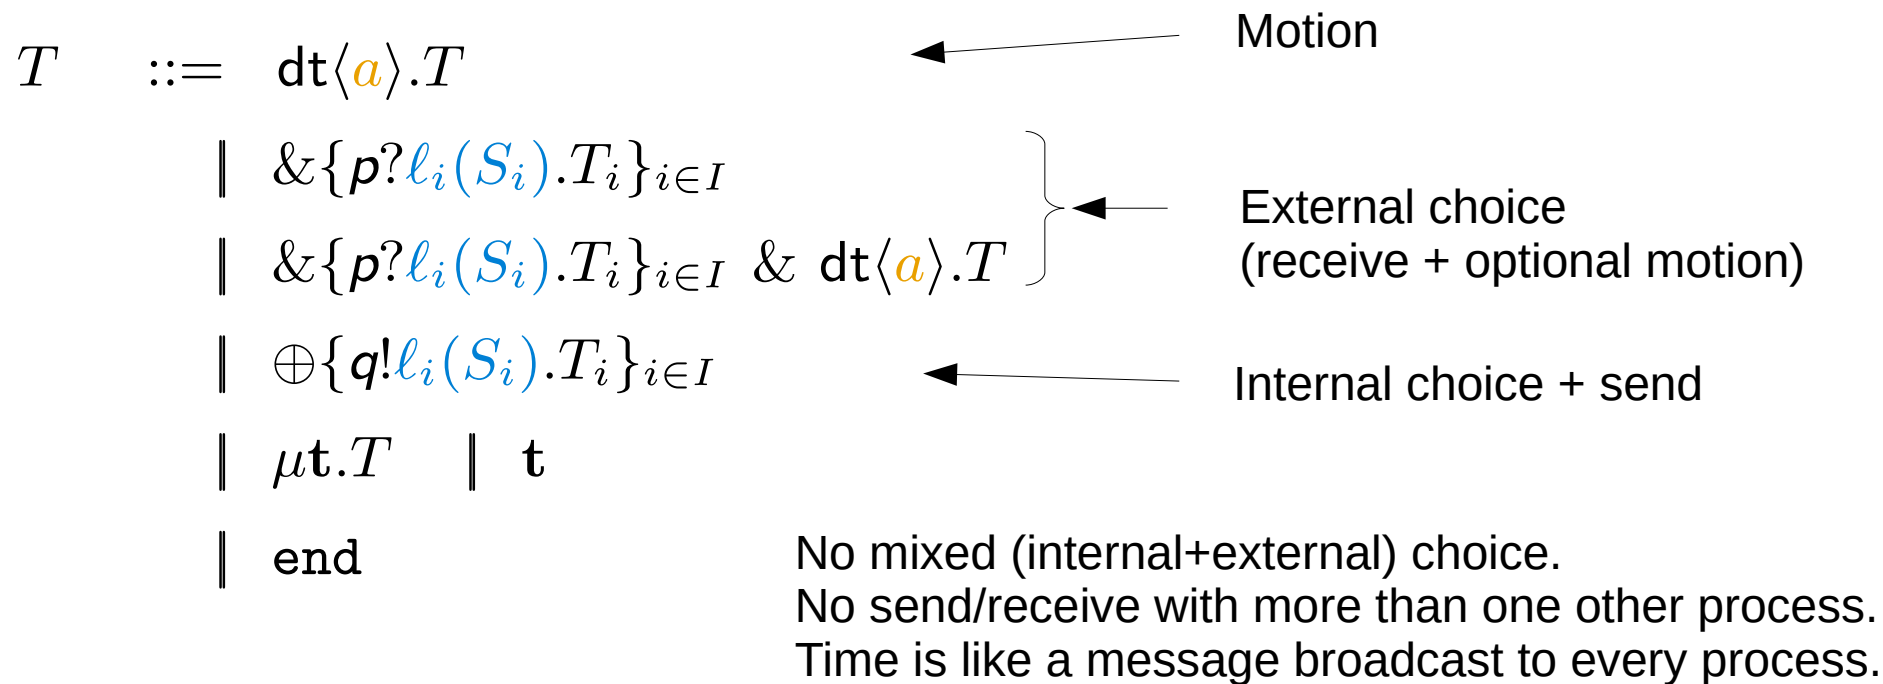
Implementation

Global Types


$G ::= \text{dt}\langle (p_i : a_i) \rangle . G$	←	Motion for each process
$\parallel p \rightarrow q : \{ \ell_i(S_i) . G_i \}_{i \in I}$	←	Communication & branching (same sender+receiver for all choices)
$\parallel \mu \mathbf{t} . G \quad \parallel \mathbf{t}$	←	Recursion (no argument \rightarrow regular language)
$\parallel \text{end}$		

Global descriptions of the messages and the motions happening in the robots.

Local Types



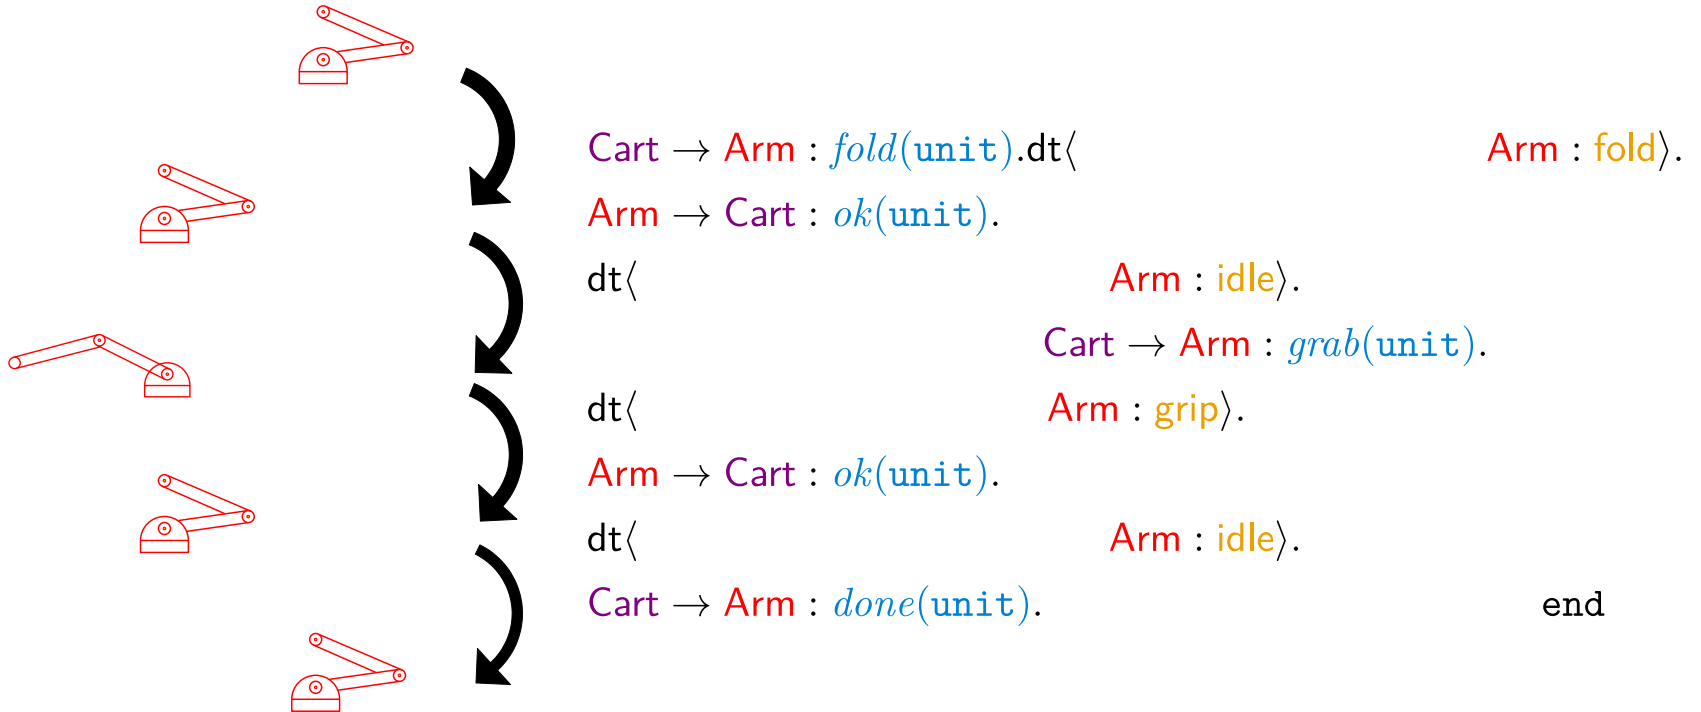
Handover: Global Type




```
Cart → Arm : fold(unit).dt⟨Cart : idle, Carrier : idle, Arm : fold⟩.  
Arm → Cart : ok(unit).Cart → Carrier : ok(unit).  
dt⟨Cart : move, Carrier : move, Arm : idle⟩.  
Carrier → Cart : ok(unit).Cart → Arm : grab(unit).  
dt⟨Cart : idle, Carrier : idle, Arm : grip⟩.  
Arm → Cart : ok(unit).Cart → Carrier : ok(unit).  
dt⟨Cart : move, Carrier : move, Arm : idle⟩.  
Cart → Arm : done(unit).Cart → Carrier : done(unit).end
```

■ Motion primitive
■ Communication

Handover: Projection on the Arm




Handover: Projection on the Arm



```
Cart?fold(unit).  
    dt<fold>.  
    Cart!ok(unit).  
dt<idle>.  
Cart?grab(unit).  
    dt<grip>.  
    Cart!ok(unit).  
dt<idle>.  
Cart?done(unit).end
```

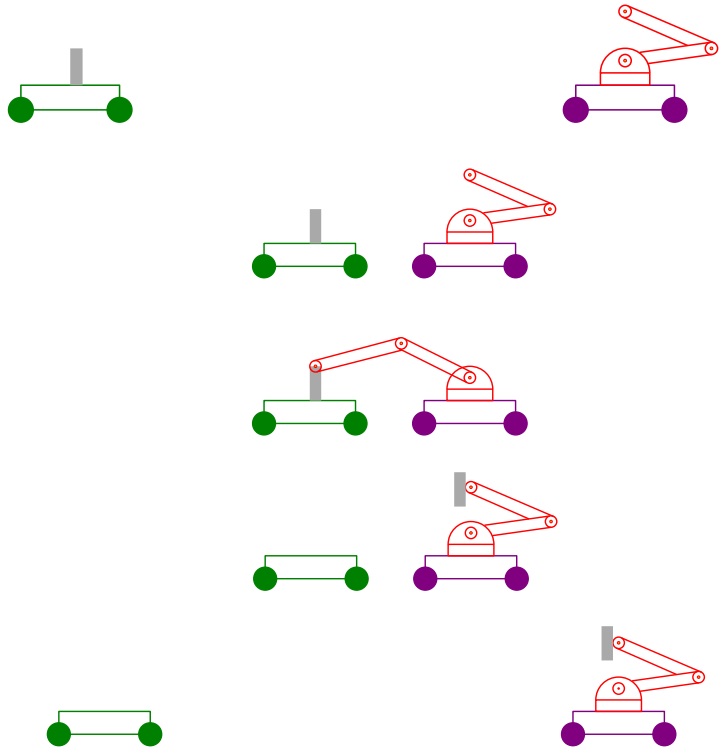
Projection is Tricky

- “Projection then product” usually adds behaviors. (geometric analogy)



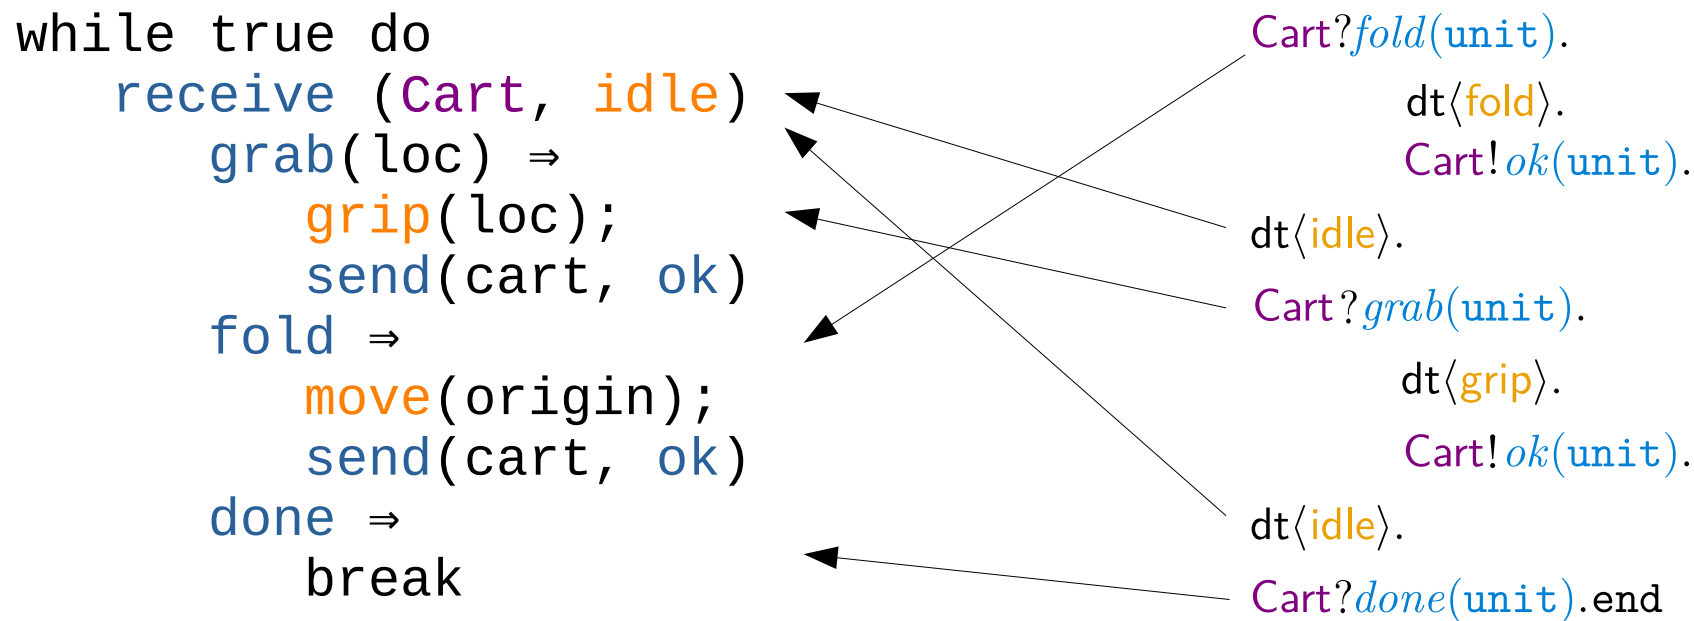
- Multiparty Session Type has a behavior preserving projection.
 - Only works on a *restricted* class of protocols.

Arm's Code



```
while true do
  receive (Cart, idle)
  grab(loc) ⇒
    grip(loc);
    send(cart, ok)
  fold ⇒
    move(origin);
    send(cart, ok)
  done ⇒
    break
```

Typing the Arm: Not So Straightforward



Subtyping: Receiving Messages

Common parts preserve subtyping.

[SUB-IN2]

$$\forall i \in I : \quad S'_i \leq S_i \quad T_i \leq T'_i$$

$$\&\{p?\ell_i(S_i).T_i\}_{i \in I} \cup J \& \text{dt}\langle a \rangle.T \leq \&\{p?\ell_i(S'_i).T'_i\}_{i \in I}$$

Can have more messages and motions.
(Never exercised)

Using Subtyping

Loop unfolded 3 times:

```
receive (Cart, idle)
  grab(loc) ⇒
    grip(loc);
    send(cart, ok)
  fold ⇒
    move(origin);
    send(cart, ok)
  done ⇒
    break
receive (Cart, idle)
  grab(loc) ⇒
    grip(loc);
    send(cart, ok)
  fold ⇒
    move(origin);
    send(cart, ok)
  done ⇒
    break
receive (Cart, idle)
  grab(loc) ⇒
    grip(loc);
    send(cart, ok)
  fold ⇒
    move(origin);
    send(cart, ok)
  done ⇒
    break
```

Cart?*fold*(unit).
dt<fold>.
Cart!*ok*(unit).


dt<idle>.

Cart?*grab*(unit).
dt<grip>.
Cart!*ok*(unit).


dt<idle>.

Cart?*done*(unit).end


Experiments




Fetch



Handover




Twist




Underpass

Language to Help Decomposition




P

Isolates Synchronization from Data



PSync

Isolates Steps in a Distributed Algorithm



PGCD

Isolates Discrete and Continuous, Components